

# 2016

## Testdokumentation



Von:  
Kathrin Donandt, Line Tayo

Goethe-Universität Frankfurt am Main  
2/21/2016

## **Inhalt**

1. Über Tajo.....	1
1.1 Was ist Tajo .....	1
1.2 Wozu Tajo .....	1
2. Installation .....	1
3. TSQL – Tajo Shell.....	2
3.1 Externe Tabellen laden.....	2
3.2 Queries ausführen .....	2
4. Hive .....	3
5. Anmerkungen .....	4

## 1. Über Tajo<sup>1</sup>

### 1.1 Was ist Tajo

Tajo ist ein opensource Data Warehouse für Apache Hadoop, das mit relationalen und verteilten Daten arbeitet. Es wird als eines der nächste Generation von Big Data Warehouse Systemen bezeichnet.

### 1.2 Wozu Tajo

Die Anwendungsgebiete von Tajo sind:

- Low-Latency und skalierbare Ad-hoc-Abfragen
- Online-Aggregation und ETL (Extraktion Transmission und Laden)
- Speicherung von große Datenmengen, wobei die Daten von HDFS oder anderen Datenquellen sein können
- direkte Kontrollmöglichkeit über verteilte Ausführungen und den Datenfluss vieler Abfrageauswertungsstrategien und Optimierungsmöglichkeiten durch SQL-Standard und erweiterte Datenbanktechniken

## 2. Installation

Tajo wurden auf der Hadoop Plattform CDH Cloudera 5.5.2 installiert. Diese Plattform wurde als Virtual Machine genutzt, wobei VMWare als Virtual Machine Environment verwendet wurde, und der Host einen Arbeitsspeicher von 8 Gigabytes besaß. Der VM wurden 4GB Memory und 2 Prozessoren zugewiesen.

Die Folgenden zwei verschiedenen Installationsanleitungen für Tajo 0.11.0 wurden getestet:

- a) Die Installationsanleitung von Apache Tajo selbst<sup>2</sup>, wobei die Konfigurationen selbst vorgenommen werden müssen. Vor der Installation von Tajo selbst müssen Hadoop (zwischen 2.3.0 und 2.6.0), Java (mindestens 1.7) und Protocol buffer 2.5.0 installiert werden. Zu beachten ist, dass die Datei: tajo-env.sh, in der die Umgebungsvariablen eingetragen werden müssen, untern „[tajo.0.11.0-src/tajo-dist-/target/tajo-0.11.0/conf](#)“ geändert werden muss.
- b) Die Installationsanleitung von Gruter<sup>3</sup>, die ein konfiguriertes Desktop-Paket mitliefert, sowie die automatische Erstellung der Tabellen und Daten der tpch Benchmark (10 MB) ermöglicht .

---

<sup>1</sup> <http://tajo.apache.org/>

<sup>2</sup> [https://tajo.apache.org/docs/current/getting\\_started.html](https://tajo.apache.org/docs/current/getting_started.html)

<sup>3</sup> <http://www.gruter.com/blog/getting-started-with-tajo-on-your-desktop/>

### 3. TSQL – Tajo Shell

In der Tajo Shell können mit einer SQL-ähnlich Syntax Datenbanken und Tabellen erstellt und mit Daten gefüllt werden. In diese Experimente (Installationsanleitung a)) würden Existierende Tabellen mit 2 GB an Daten in Tajo geladen. Diese Daten waren zuvor nach der Anleitung unter <https://github.com/t-ivanov/D2F-Bench> generiert worden. Die Tajo Shell startet man mit dem Befehl ./tsql (sudo ./tsql) unter `tajo.0.11.0-src/tajo-dist-/target/tajo-0.11.0/bin`.

#### 3.1 Externe Tabellen laden

Optional: Man kann selbst eine Datenbank erstellen (CREATE DATABASE xyz).

Die Syntax zum Laden einer externen Tabelle lautet wie folgt (hier am Beispiel der Tabelle „NATION“ die unter `home/cloudera/[...]/nation.tbl` lokal vorlag):

```
create external table NATION (
    N_NATIONKEY INTEGER [NOT NULL],
    N_NAME CHAR(25),
    N_REGIONKEY INTEGER [NOT NULL],
    N_COMMENT [VARCHAR(152)] TEXT)
using text with ('text.delimiter'='|') location 'file:/home/cloudera/[...]/nation.tbl';
```

Zu beachten ist, dass Tajo **nicht alle SQL Ausdrücke** unterstützt.

Liegt die Tabelle auf HDFS vor, muss man statt „file :/“ „hdfs :/“ verwenden.

#### 3.2 Queries ausführen

Bei der Ausführung von Queries sind wieder Syntaxunterschiede zu SQL zu beachten. Beispielweise muss ein Datum explizit als Datei geparsst werden:

```
select l_returnflag,l_linenumber,...from lineitem
where l_shipdate <= '1998-09-01' :: date
group by l_returnflag, l_linenumber
order by l_returnflag, l_linenumber;
```

Die Status der Queries, kann man sich unter „localhost:26080/query.jsp“ anzeigen lassen:

The screenshot shows the Tajo Master interface at [localhost:26080/query.jsp](http://localhost:26080/query.jsp). It displays three main sections: Submitted Queries, Running Queries, and Finished Queries.

- Submitted Queries:** Shows two queries:
  - q\_1454104630983\_0001: Submitted on 2016-01-29 13:58:02, Progress 0%, Status QUERY\_MASTER\_INIT. SQL: select ... (long query)
  - q\_1454104630983\_0003: Submitted on 2016-01-31 08:31:33, Progress 0%, Status QUERY\_MASTER\_INIT. SQL: select ... (long query)
- Running Queries:** No running queries.
- Finished Queries:** One query listed:
 

QueryId	Query Master	Started	Finished	Time	Status	sql
q_1454104630983_0003	null	2016-01-31 08:25:39	2016-01-31 08:25:39	15 msec	QUERY_SUCCEEDED	select * from tpc_h10m.nation

Wie auf dem Bild zu sehen ist, haben komplexe Queries nie den Status „Running“ erreicht, sondern sind im „Submitted“ Status hängen geblieben.

Nach Testen der Hive Integration haben die komplexen Queries dann den Status „Running“ und „Finished“ erreicht, aber die Ergebnistabellen waren leer bzw. es wurde kein Ergebniswert geliefert.

## 4. Hive

Die ursprüngliche Absicht war es, Tajo mit Hive zu vergleichen. Die Anleitung für Hive-testbench findet man unter [github.com/t-ivanov/hive-testbench](https://github.com/t-ivanov/hive-testbench) und für 2 GB Daten erhielten wir folgende Ergebnisse:

Query	Hive	Query	Hive
1	139.877 s	12	286.675 s
2	78.58 s	13	391.11 s
3	293.375 s	14	100.994 s
4	284.408 s	15	366.933 s
5	363.678 s	16	122.3 s
6	133.737 s	17	497.794 s
7	760.074 s (12 min)	18	893.839 s

<b>8</b>	360.757 s	<b>19</b>	208.998 s
<b>9</b>	1382.928 s (23 min)	<b>20</b>	83.396 s
<b>10</b>	346.588 s	<b>21</b>	1057.563 s
<b>11</b>	406.386 s	<b>22</b>	345.178 s

## 5. Anmerkungen

Bei der Arbeit mit der Tajo Shell wurden folgende Unterschiede zu SQL festgestellt.

- Sie akzeptiert auch keine verschachtelten select Anfragen.
- Tajo akzeptiert keine *PRIMARY KEYS*
- „*NOT NULL*“ (leeren Zellen) sind nicht erlaubt
- keine Views
- *CHAR* ok, *VARCHAR* nicht akzeptiert (stattdessen: *TEXT*) ,
- *DEZIMAL* nicht erlaubt (stattdessen *FLOAT*)
- case-insensitive

Als Fazit kann man festhalten, dass die Arbeit mit Tajo nicht straight-forward war, und unsere Hardwarekapazitäten für das Testen nicht optimal waren.

## Documentation

# Tajo Benchmark on Cloudera CDH

Andreas Grimm



February 2016



Goethe-Universität Frankfurt am Main  
Fachbereich 12: Informatik und Mathematik  
Institut für Informatik

**Working time:**  
Oktober 2015 - February 2016

**Supervisor:**  
Todor Ivanov

## Contents

<b>1 Introduction</b> . . . . .	1
<b>2 Installation</b> . . . . .	1
2.1 Prerequisites . . . . .	1
2.2 Download tajo . . . . .	2
2.3 Configure the project pom file with the CDH5 maven repository . . . . .	2
<b>3 Configuration</b> . . . . .	2
<b>4 TPCH Benchmark</b> . . . . .	3
<b>5 Results</b> . . . . .	3
<b>A Summary</b> . . . . .	4
<b>B Modifications on pom files</b> . . . . .	6
B.1 tajo-project . . . . .	6
B.2 tajo-catalog . . . . .	6
B.3 tajo-hcatalog . . . . .	7
<b>C Import statements</b> . . . . .	12
C.1 Supplier . . . . .	12
C.2 Lineitem . . . . .	12
C.3 Part . . . . .	12
C.4 Partsupp . . . . .	13
C.5 Customer . . . . .	13
C.6 Orders . . . . .	13
C.7 Nation . . . . .	14
C.8 Region . . . . .	14

# 1 Introduction

About Tajo:

- compatible with ANSI/ISO SQL - SQL like syntax
- advanced query optimization and lot's of optimization techniques possible
- used for low-latency and scalable ad-hoc queries
- HDFS or external sources for data possible
- user defined functions, interactive shell

# 2 Installation

## 2.1 Prerequisites

- i) Hadoop 2.3.0 or higher (up to 2.6.0)
- ii) Java 1.7 or higher
- iii) Protocol buffer 2.5.0

The first two work out of the box on CDH 5.4.2 but the protocol buffer 2.5.0 has to be installed manually. The steps are:

```
>> wget http://protobuf.googlecode.com/files/protobuf-2.5.0.tar.gz
>> tar -zxf protobuf-2.5.0.tar.gz
>> cd protobuf-2.5.0
>> ./configure
>> make
>> sudo make install
>> export PATH=$PATH:/usr/local/bin
```

It is recommended to install the commons-logging-1.2 package manually. Go to [http://commons.apache.org/proper/commons-logging/download\\_logging.cgi](http://commons.apache.org/proper/commons-logging/download_logging.cgi) and download the source package (commons-logging-1.2-src.tar.gz). Go to the target directory of the download and do the following:

```
>> cd commons-logging-1.2-src
>> mvn package
```

If the build fails, try to add the following plugin to the pom file:

```
<plugin>
    <groupId>org.codehaus.mojo</groupId>
    <artifactId>exec-maven-plugin</artifactId>
    <version>1.2</version>
</plugin>
```

## 2.2 Download tajo

Try to get tajo using

```
>> git clone https://git-wip-us.apache.org/repos/asf/tajo.git
```

This failed in my case. The other way, which is also recommended, is to download it directly from <http://www.apache.org/dyn/closer.cgi/tajo>. I used tajo version 0.10.0. This makes sense, since the tutorials that can be found out in the net are mostly for 0.9.0 or 0.8.0. Using a even higher version than 0.10.0 might cause unnecessary troubles due to incompatibilities. Navigate to the download directory:

```
>> tar xzvf tajo-0.10.0-src.tar.gz
```

## 2.3 Configure the project pom file with the CDH5 maven repository

Do some modifications on several pom.xml files as described at <https://cwiki.apache.org/confluence/display/TAJO/How+to+Build+Apache+Tajo+for+CDH5>. The modifications for cloudera CDH 5.4.2 can be found in Appendix B. Note that some of these mods are not necessary for the following. After the modification go to tajo-0.10.0-src (in my case /home/cloudera/Desktop/tajo-0.10.0-src/) and execute

```
>> mvn clean install package -DskipTests -Pdist  
-Dhadoop.version=2.6.0-cdh5.4.2
```

This should create a local build of apache tajo on CDH 5.4.2 on CentOs 6.4 using Red Hat as operating system.

## 3 Configuration

Now do the required modifications on the config file. Since I downloaded the tajo distribution to the desktop the path is:

/home/cloudera/Desktop/tajo-0.10.0-src/tajo-dist/target/tajo-0.10.0/ Open /conf/tajo-env.sh and add the following lines:

```
export /usr/lib/hadoop  
export JAVA_HOME=/usr/java/jdk1.7.0_67-cloudera  
export TAJO_CLASSPATH=$HADOOP_HOME/client/*:$HADOOP_HOME/lib/*
```

Some sources say that the hadoop home is /opt/cloudera/parcels/CDH/lib/hadoop but that is only the case if you upgraded using parcels. The tutorial on the official tajo homepage says that the last path is not needed, but in my case it was to run tajo. After that the local tajo cluster should be ready to start:

```
sudo ./bin/start-tajo.sh
```

Finally you can start the tajo-shell:

```
sudo ./bin/tsql
```

After each session one should probably stop the cluster using:

```
sudo ./bin/stop-tajo.sh
```

## 4 TPCH Benchmark

First create a new database:

```
CREATE DATABASE tpch;
```

```
\c tpch;
```

Import tables and data using the statements listed in App. C. The external data was created using the data generator from <https://github.com/xinglin/tpch> and the step-by-step guide from <https://sites.google.com/site/halitsch88/Implementation-TPC-H-schema-into-MySQL-DBMS>.

Most of the TPCH queries have to be modified. The following sources were used:

<https://cwiki.apache.org/confluence/display/TAJO/TPC-H+Benchmark+Set>

[https://s3-us-west-2.amazonaws.com/tajo/taas/documents/Analyze\\_Data\\_with\\_Tajo.pdf](https://s3-us-west-2.amazonaws.com/tajo/taas/documents/Analyze_Data_with_Tajo.pdf)

Since some queries have to be split into several commands it's useful to measure the total execution time via

```
/usr/bin/time -f Time:%e ./tsql -f querieX.sql databaseY > resultZ.txt
```

where querieX.sql is the executed query, databaseY the database with the stored data, and the results are stored in resultZ.txt.

## 5 Results

Results for the benchmarks using 1GB in the case innodb vs. tajo and 2GB in the case hive vs tajo.

Query	Tajo	InnoDB	Query	Tajo	InnoDB
1	26.49 s	43.689116 s	7	37.81 s	9.813247 s
2	10.48 s	3.123255 s	8	25.23 s	15.086646 s
3	32.13 s	67.431433 s	9	47.95 s	137.599134 s
4	51.24 s	3.862811 s	10	20.15 s	55.171110 s
5	2253.99 s	6.768237 s	11	9.12 s	1.450574 s
6	21.28 s	4.112015 s	12	43.43 s	10.804388 s

Query	Tajo	InnoDB	Query	Tajo	InnoDB
13	13.04 s	46.216578 s	18	30.78 s	27552.101274 s
14	21.38 s	406.764523 s	19	6391.68 s	1.407827 s
15	23.57 s	7.939032 s	20	15.75 s	2.537518 s
16	11.88 s	1.306451 s	21	81.85 s	24.867429 s
17	30.82 s	1.497502 s	22	5.45 s	0.610497 s

Query	Tajo	Hive	Query	Tajo	Hive
1	26.87 s	139.877 s	7	36.93 s	760.074 s (12 min)
2	6.87 s	78.58 s	8	23.88 s	360.757 s
3	28.51 s	293.375 s	9	62.06 s	1382.928 s (23 min)
4	41.29 s	284.408 s	10	14.11 s	346.588 s
5	7535.25 s	363.678 s	11	4.74 s	406.386 s
6	19.97 s	133.737 s	12	42.54 s	286.675 s

Query	Tajo	Hive	Query	Tajo	Hive
13	11.45 s	391.11 s	18	34.55 s	893.839 s
14	19.28 s	100.994 s	19		208.998 s
15	20.11 s	366.933 s	20	28.13 s	83.396 s
16	9.21 s	122.3 s	21	357.79 s	1057.563 s
17	30.77 s	497.794 s	22	7.50 s	345.178 s

## A Summary

- no primary keys supported by tajo
- modified sql syntax: ‘10-10-2010’::date
- no scalar subqueries
- no views
- case-insensitive
- “not null“ not allowed
- use text in stead of varchar

- buggy error messages: Error: no error message
- some queries hang without reason
- drop table if exists sometimes does not seem to work, temporary tables had to be dropped manually

## B Modifications on pom files

All paths in this section refer to the directory you unpacked the tajo distribution. For example in my case /tajo-project/pom.xml referred to the directory /home/cloudera/Desktop/tajo-0.10.0-src/tajo-project/pom.xml assuming you downloaded the tajo distribution to the desktop.

### B.1 tajo-project

If one of the following is missing in /tajo-project/pom.xml, add it.

```
<repositories>
    <repository>
        <id>apache . snapshots </id>
        <url>http :// repository . apache . org / snapshots </url>
        <snapshots>
            <enabled>true </enabled>
        </snapshots>
    </repository >

    <repository>
        <id>eclipse - jetty </id>
        <url>
            http :// repo2 . maven . org / maven2 / org / eclipse / jetty / jetty - distribution /
        </url>
        <snapshots>
            <enabled>false </enabled>
        </snapshots>
    </repository >

    <repository>
        <id>cloudera </id>
        <url>
            https :// repository . cloudera . com / artifactory / cloudera - repos /
        </url>
        <snapshots>
            <enabled>false </enabled>
        </snapshots>
    </repository >
</repositories >
```

### B.2 tajo-catalog

Add the following profile to tajo's catalog pom file at /tajo-catalog/tajo-catalog-drivers/pom.xml:

```

<profile>
    <id>hcatalog-cdh5.4.2</id>
    <activation>
        <activeByDefault>false</activeByDefault>
    </activation>
    <modules>
        <module>tajo-hcatalog</module>
    </modules>
</profile>

```

Adding the following to /tajo-dist/pom.xml did not work for me:

```

if [ -f $ROOT/tajo-catalog/tajo-catalog-drivers/tajo-hcatalog/target/lib/
      hive-hcatalog-core--cdh*.jar ]
then
run cp -r $ROOT/tajo-catalog/tajo-catalog-drivers/tajo-hcatalog/target/
      lib/hive-hcatalog-core--cdh*.jar lib/
fi

```

### B.3 tajo-hcatalog

/tajo-catalog/tajo-catalog-drivers/tajo-hcatalog/pom.xml:

```

<profile>
    <repositories>
        <repository>
            <id>cloudera</id>
            <url>
                https://repository.cloudera.com/artifactory/cloudera-repos/
            </url>
            <snapshots>
                <enabled>true</enabled>
            </snapshots>
        </repository>
    </repositories>
    <id>hcatalog-cdh5.4.2</id>
    <activation>
        <activeByDefault>false</activeByDefault>
    </activation>
    <properties>
        <hive.version>1.1.0-cdh5.4.2</hive.version>
    </properties>
    <dependencies>
        <dependency>
            <groupId>javax.jdo</groupId>

```

```

<artifactId>jdo2-api </artifactId>
<version>2.3-eb</version>
<scope>provided </scope>
</dependency>
<dependency>
    <groupId>org.apache.hive </groupId>
    <artifactId>hive-exec </artifactId>
    <version>${hive.version} </version>
    <scope>provided </scope>
    <exclusions>
        <exclusion>
            <groupId>org.apache.hive </groupId>
            <artifactId>hive-common </artifactId>
        </exclusion>
        <exclusion>
            <groupId>org.apache.hive </groupId>
            <artifactId>hive-contrib </artifactId>
        </exclusion>
        <exclusion>
            <groupId>org.apache.hive </groupId>
            <artifactId>hive-hbase-handler </artifactId>
        </exclusion>
        <exclusion>
            <groupId>org.apache.hive </groupId>
            <artifactId>hive-metastore </artifactId>
        </exclusion>
        <exclusion>
            <groupId>org.apache.hive </groupId>
            <artifactId>hive-serde </artifactId>
        </exclusion>
        <exclusion>
            <groupId>org.apache.hive </groupId>
            <artifactId>hive-shims </artifactId>
        </exclusion>
        <exclusion>
            <groupId>org.apache.hive </groupId>
            <artifactId>hive-testutils </artifactId>
        </exclusion>
        <exclusion>
            <groupId>org.apache.thrift </groupId>
            <artifactId>libfb303 </artifactId>
        </exclusion>
        <exclusion>
            <groupId>org.apache.thrift </groupId>

```

```

        <artifactId>libthrift </artifactId>
    </exclusion>
    <exclusion>
        <groupId>javax.jdo </groupId>
        <artifactId>jdo2-api </artifactId>
    </exclusion>
</exclusions>
</dependency>
<dependency>
    <groupId>org.apache.hive </groupId>
    <artifactId>hive-metastore </artifactId>
    <version>${hive.version} </version>
    <scope>provided </scope>
    <exclusions>
        <exclusion>
            <groupId>org.apache.hive </groupId>
            <artifactId>hive-common </artifactId>
        </exclusion>
        <exclusion>
            <groupId>org.apache.hive </groupId>
            <artifactId>hive-serde </artifactId>
        </exclusion>
        <exclusion>
            <groupId>org.apache.hive </groupId>
            <artifactId>hive-shims </artifactId>
        </exclusion>
        <exclusion>
            <groupId>org.apache.thrift </groupId>
            <artifactId>libfb303 </artifactId>
        </exclusion>
        <exclusion>
            <groupId>org.apache.thrift </groupId>
            <artifactId>libthrift </artifactId>
        </exclusion>
        <exclusion>
            <groupId>javax.jdo </groupId>
            <artifactId>jdo2-api </artifactId>
        </exclusion>
    </exclusions>
</dependency>
<dependency>
    <groupId>org.apache.hive </groupId>
    <artifactId>hive-cli </artifactId>
    <version>${hive.version} </version>

```

```

<scope>provided </scope>
<exclusions>
    <exclusion>
        <groupId>org . apache . hive </groupId>
        <artifactId>hive-common</artifactId>
    </exclusion>
    <exclusion>
        <groupId>org . apache . hive </groupId>
        <artifactId>hive-exec </artifactId>
    </exclusion>
    <exclusion>
        <groupId>org . apache . hive </groupId>
        <artifactId>hive-metastore </artifactId>
    </exclusion>
    <exclusion>
        <groupId>org . apache . hive </groupId>
        <artifactId>hive-serde </artifactId>
    </exclusion>
    <exclusion>
        <groupId>org . apache . hive </groupId>
        <artifactId>hive-service </artifactId>
    </exclusion>
    <exclusion>
        <groupId>org . apache . hive </groupId>
        <artifactId>hive-shims </artifactId>
    </exclusion>
    <exclusion>
        <groupId>javax . jdo </groupId>
        <artifactId>jdo2-api </artifactId>
    </exclusion>
</exclusions>
</dependency>
<dependency>
    <groupId>org . apache . hive . hcatalog </groupId>
    <artifactId>hive-hcatalog-core </artifactId>
    <version>${hive . version}</version>
    <scope>provided </scope>
    <exclusions>
        <exclusion>
            <groupId>org . apache . hive </groupId>
            <artifactId>hive-cli </artifactId>
        </exclusion>
        <exclusion>
            <groupId>org . apache . hive </groupId>

```

```

        <artifactId>hive-common</artifactId>
    </exclusion>
    <exclusion>
        <groupId>org.apache.hive</groupId>
        <artifactId>hive-exec</artifactId>
    </exclusion>
    <exclusion>
        <groupId>org.apache.hive</groupId>
        <artifactId>hive-metastore</artifactId>
    </exclusion>
    <exclusion>
        <groupId>org.apache.hive</groupId>
        <artifactId>hive-serde</artifactId>
    </exclusion>
    <exclusion>
        <groupId>org.apache.hive</groupId>
        <artifactId>hive-service</artifactId>
    </exclusion>
    <exclusion>
        <groupId>org.apache.hive</groupId>
        <artifactId>hive-shims</artifactId>
    </exclusion>
    <exclusion>
        <groupId>javax.jdo</groupId>
        <artifactId>jdo2-api</artifactId>
    </exclusion>
    </exclusions>
</dependency>
<dependency>
    <groupId>org.apache.hadoop</groupId>
    <artifactId>hadoop-mapreduce-client-core</artifactId>
    <version>${hadoop.version}</version>
    <scope>provided</scope>
</dependency>
</dependencies>
</profile>

```

## C Import statements

### C.1 Supplier

```
create external table SUPPLIER (
    S_SUPPKEY bigint ,
    S_NAME text ,
    S_ADDRESS text ,
    S_NATIONKEY bigint ,
    S_PHONE text ,
    S_ACCTBAL double ,
    S_COMMENT text)
using text with ('text.delimiter'='|')
location 'file :/ path/to/table/supplier.tbl';
```

### C.2 Lineitem

```
create external table LINEITEM (
    L_ORDERKEY bigint ,
    L_PARTKEY bigint ,
    L_SUPPKEY bigint ,
    L_LINENUMBER bigint ,
    L_QUANTITY double ,
    L_EXTENDEDPRICE double ,
    L_DISCOUNT double ,
    L_TAX double ,
    L_RETURNFLAG text ,
    L_LINESTATUS text ,
    L_SHIPDATE date ,
    L_COMMITDATE date ,
    L_RECEIPTDATE date ,
    L_SHIPINSTRUCT text ,
    L_SHIPMODE text ,
    L_COMMENT text)
using text with ('text.delimiter'='|')
location 'file :/ path/to/table/lineitem.tbl';
```

### C.3 Part

```
create external table PART (
    P_PARTKEY bigint ,
    P_NAME text ,
    P_MFGR text ,
    P_BRAND text ,
    P_TYPE text ,
```

```

P_SIZE integer ,
P_CONTAINER text ,
P_RETAILPRICE double ,
P_COMMENT text)
using text with ('text.delimiter'='|')
location 'file :/ path/to/table/part.tbl';

```

#### C.4 Partsupp

```

create external table PARTSUPP (
    PS_PARTKEY bigint ,
    PS_SUPPKEY bigint ,
    PS_AVAILQTY int ,
    PS_SUPPLYCOST double ,
    PS_COMMENT text)
using text with ('text.delimiter'='|')
location 'file :/ path/to/table/partsupp.tbl';

```

#### C.5 Customer

```

create external table CUSTOMER (
    C_CUSTKEY bigint ,
    C_NAME text ,
    C_ADDRESS text ,
    C_NATIONKEY bigint ,
    C_PHONE text ,
    C_ACCTBAL double ,
    C_MKTSEGMENT text ,
    C_COMMENT text)
using text with ('text.delimiter'='|')
location 'file :/ path/to/table/customer.tbl';

```

#### C.6 Orders

```

create external table ORDERS (
    O_ORDERKEY bigint ,
    O_CUSTKEY bigint ,
    O_ORDERSTATUS text ,
    O_TOTALPRICE double ,
    O_ORDERDATE date ,
    O_ORDERPRIORITY text ,
    O_CLERK text ,
    O_SHIPPRIORITY int ,
    O_COMMENT text)
using text with ('text.delimiter'='|')
location 'file :/ path/to/table/orders.tbl';

```

## C.7 Nation

```
create external table NATION (
    N_NATIONKEY bigint,
    N_NAME text,
    N_REGIONKEY bigint,
    N_COMMENT text)
using text with ('text.delimiter'='|')
location 'file:/path/to/table/nation.tbl';
```

## C.8 Region

```
create external table REGION (
    R_REGIONKEY bigint,
    R_NAME text,
    R_COMMENT text)
using text with ('text.delimiter'='|')
location 'file:/path/to/table/region.tbl';
```